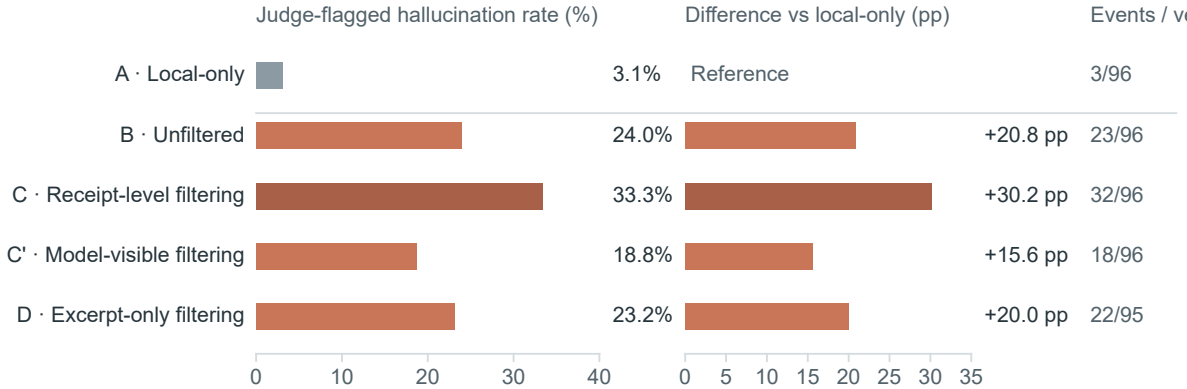
Judge-flagged hallucination rate (%)
Difference vs local-only (pp)
Events / verdicts
A · Local-only
3.1%
Reference
3/96
B · Unfiltered
24.0%
+20.8 pp
23/96
C · Receipt-level filtering
33.3%
+30.2 pp
32/96
C' · Model-visible filtering
18.8%
+15.6 pp
18/96
D · Excerpt-only filtering
23.2%
+20.0 pp
22/95
0
10
20
30
40
0
5
10
15
20
25
30
35
Old architecture (pre-separation) · 2026-09-23/24
48-question subset · two independent judges · 96 verdicts per arm (differences are descriptive)